

Introduction to Machine Learning (67577)

Exercise 6 Convex Optimization, Gradient Descent

Second Semester, 2022

Contents

1	Submission Instructions	2
2	Theoretical Part	2
2.1	Convex optimization	2
2.2	Sub-gradients for Soft-SVM Objective	2
3	Practical Part	3
3.1	Gradient Descent	3
3.1.1	Comparing Fixed learning rates	4
3.1.2	Comparing Exponentially Decaying learning rates	4
3.2	Minimizing Regularized Logistic Regression	4

1 Submission Instructions

Please make sure to follow the general submission instructions available on the course website. In addition, for the following assignment, submit a single `ex6_ID.tar` file containing:

- An `Answers.pdf` file with the answers for all theoretical and practical questions (include plotted graphs *in* the PDF file).
- The following python files (without any directories): `gradient_descent.py`, `learning_rate.py`, `modules.py`, `logistic_regression.py` and `gradient_descent_investigation.py`

The `ex6_ID.tar` file must be submitted in the designated Moodle activity prior to the date specified *in the activity*.

- Late submissions will not be accepted and result in a zero mark.
- Plots included as separate files will be considered as not provided.
- Do not forget to answer the Moodle quiz of this assignment.

2 Theoretical Part

2.1 Convex optimization

1. Let $f_1, \dots, f_m : C \rightarrow \mathbb{R}$ be a set of convex functions and $\gamma_1, \dots, \gamma_m \in \mathbb{R}_+$. Prove from definition that $g(\mathbf{u}) = \sum_{i=1}^m \gamma_i f_i(\mathbf{u})$ is a convex function.
2. Give a counterexample for the following claim: Given two functions $f, g : \mathbb{R} \rightarrow \mathbb{R}$, define a new function $h : \mathbb{R} \rightarrow \mathbb{R}$ by $h = f \circ g$. If f and g are convex then h is convex as well.
3. Let $f : C \rightarrow \mathbb{R}$ be a function defined over a convex set C . Prove that f is convex iff its *epigraph* is a convex set, where $\text{epi}(f) = \{(u, t) : f(u) \leq t\}$.
4. Let $f_i : V \rightarrow \mathbb{R}, i \in I$. Let $f : V \rightarrow \mathbb{R}$ given by

$$f(u) = \sup_{i \in I} f_i(u).$$

If f_i are convex for every $i \in I$, then f is also convex.

2.2 Sub-gradients for Soft-SVM Objective

The Soft-SVM objective, though convex, is not differentiable in all of its domain due to the use of the hinge-loss. Therefore, to implement a sub-gradient descent solver for this problem we must first describe sub-gradients of the objective.

5. Given $\mathbf{x} \in \mathbb{R}^d$ and $y \in \{\pm 1\}$. Show that the hinge loss is convex in \mathbf{w}, b . That is, define

$$f(\mathbf{w}, b) := \ell_{\mathbf{x}, y}^{\text{hinge}}(\mathbf{w}, b) = \max(0, 1 - y(\mathbf{x}^\top \mathbf{w} + b))$$

and show that f is convex in \mathbf{w}, b .

6. Deduce some sub-gradient of the hinge loss function $g \in \partial \ell_{\mathbf{x}, y}^{\text{hinge}}(\mathbf{w}, b)$.
7. Let $f_1, \dots, f_m : \mathbb{R}^d \rightarrow \mathbb{R}$ be a set of convex functions and $\mathbf{g}_k \in \partial f_k(\mathbf{x})$ for all $k \in [m]$ be sub-gradients of these functions. Define $f : \mathbb{R}^d \rightarrow \mathbb{R}$ by $f(\mathbf{x}) = \sum_{i=1}^m f_i(\mathbf{x})$. Show that $\sum_k \mathbf{g}_k \in \partial \sum_k f_k(\mathbf{x})$.

8. Let $S = \{(\mathbf{x}_i, y_i)\}_{i=1}^m \subseteq \mathbb{R}^d \times \{\pm 1\}$ be a sample and define $f: \mathbb{R}^d \rightarrow \mathbb{R}$ by:

$$f(\mathbf{w}, b) = \frac{1}{m} \sum_{i=1}^m \ell_{\mathbf{x}_i, y_i}^{\text{hinge}}(\mathbf{w}, b) + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

Find a sub-gradient of f for any \mathbf{w} .

3 Practical Part

Before starting the practical part please make sure to have cloned the IML.HUJI GitHub repository and setup the virtual environment as specified in the instructions. Write the necessary code in the files specified in the questions.

3.1 Gradient Descent

In the following section you will implement a generic Gradient Descent algorithm, and explore and visualize its performance on different objective functions. To assist you with the implementation please start by reading the documentation of the `GradientDescent` class in the `IMLearn.descent_methods.gradient_descent.py` file and following the steps as described below.

Learning Rate: The `GradientDescent` class, when initialized receives a *learning rate strategy* in the form of a `BaseLR` instance. Read the documentation of the `BaseLR` base class in the `IMLearn.base.base_learning_rate.py` file and then implement two learning rate strategies in the :

- Constant (Fixed) Learning Rate (i.e. $\eta_t = \eta$) - `FixedLR` class in the `IMLearn.descent_methods.learning_rate.py` file
- Exponentially decaying Learning Rate. $\eta_t = \eta \cdot \gamma^t$ - `ExponentialLR` class in the `IMLearn.descent_methods.learning_rate.py` file

Objective Functions (Modules): When running the `GradientDescent.fit` function it receives an instance derived from the `BaseModule` class. This class defines the generic abstract form of any objective to be minimized using gradient descent. Its two main functions are used to compute the value of the function and the derivative of the function at a given point of interest. Read the documentation of the `BaseModule` base class in the `IMLearn.base.base_module.py` file.

Implement the L2 and L1 modules in the `IMLearn.descent_methods.modules.py` file. Note that both these modules ignore any passed inputs in the `compute_output` and `compute_jacobian` functions and simply use the `weights` defined in the base class. Other modules implemented later (and in the next exercise will be using such inputs)

Gradient Descent Algorithm Implement the `GradientDescent` class in the `descent_methods.gradient_descent.py` file.

- Note that when instantiating a `GradientDescent` object a `callback` can be passed. This will be used to investigate different properties of the algorithm's run and will be specified in the questions below.

- Implementation must support several solution types, one of which is the average of $\mathbf{w}^{(1)}, \dots, \mathbf{w}^{(t)}$. In your implementation do not store all solutions to avoid wasting memory.

3.1.1 Comparing Fixed learning rates

We begin with investigating the GD convergence over the L1 and L2 objectives using fixed learning rates. In the `gradient_descent_investigation.py` file implement the `compare_fixed_learning_rates` function as specified in function documentation:

- Implement the `get_gd_state_recorder_callback` function as specified in function documentation. This function returns a “fresh” callback function and lists for losses and weights throughout the GD iterations.
- Minimize the L1 and L2 modules for each of the following fixed learning rates $\eta \in \{1, 0.1, 0.01, 0.001\}$, setting the initial starting point (i.e. the initial value of the module’s weights) to $\mathbf{x}_0 = \left(\sqrt{2}, \frac{e}{3}\right)$

Then, answer the following questions:

1. Plot the descent trajectory for each of the settings described above (you can use the `plot_descent_path`). Add below the plots for $\eta = 0.01$ and explain the differences seen between the L1 and L2 modules.
2. Describe two phenomena that can be seen in the descent path of the ℓ_1 objective when using GD and a fixed learning rate.
3. For each of the modules, plot the convergence rate (i.e. the norm as a function of the GD iteration) for all specified learning rates. Explain your results
4. What is the lowest loss achieved when minimizing each of the modules? Explain the differences

3.1.2 Comparing Exponentially Decaying learning rates

Next, we will use the exponential decay (instead of the fixed learning rate) to optimize the L1 module. Starting from $\mathbf{x}_0 = \left(\sqrt{2}, \frac{e}{3}\right)$ use the exponential decay with $\eta = 0.1$ and $\gamma \in \{0.9, 0.95, 0.99, 1\}$. Then, answer the following questions.

5. Plot the convergence rate for all decay rates in a single plot. Explain your results.
6. How does the algorithm perform using the exponential decay compared to the fixed learning rate? What is the lowest ℓ_1 norm achieved using the exponential decay. Explain why there are differences.
7. Plot the 2D trajectory of both Modules (using `go.Contour` and `go.Scatter`). Do this only for the $\gamma = 0.95$. Describe how the trajectory changed from when using a fixed learning rate.

3.2 Minimizing Regularized Logistic Regression

In the following part you will minimize a regularized (ℓ_1 or ℓ_2) logistic regression classifier. Implement the following as described below:

- Implement the `LogisticModule` in the `IMLearn.descent_methods.modules.py` file, as described in class documentation.

- In the `compute_output`, you should return the negative log-likelihood

$$f(\mathbf{w}) = -\log \left(\prod_i P(Y = y_i | X = \mathbf{x}_i, \mathbf{w}) \right)$$

- In the `compute_jacobian`, you should return the derivative of the objective above $\frac{\partial f(\mathbf{w})}{\partial \mathbf{w}}$.
- Recall the calculations done in recitations 5 and 11.
- Implement the `RegularizedModule` in the `IMLearn.descent_methods.modules.py` file, as described in class documentation. This module receives two generic modules to be used as the fidelity and regularization terms. For example: `RegularizedModule(LogisticModule(), L1())`.
- Implement the `LogisticRegression` class in the `IMLearn.learnres.classifiers.logistic_regression.py` file as specified in class documentation. This class should wrap the usage of your gradient descent implementation on the `LogisticModule`, `RegularizedModule`, `L1` and `L2`.

R Notice that the `LogisticRegression` class does not create an instance of the `GradientDescent` class. Instead, it receives it as a *dependency* in the constructor. As such, your `LogisticRegression` implementation is open for future extensions and support of different solvers. This is one of the 5 **SOLID** coding principles - if you wish to write good code and have a good design - be **SOLID**.

Then, load the South Africa Heart Disease dataset (`SAheart.data`), split it to train- and test sets (80% train) and answer the following questions:

8. Using your implementation, fit a logistic regression model over the data. Use the `predict_proba` to plot an ROC curve (for $\alpha \in \{0, .01, .02, \dots, .99, 1\}$).
9. Which value of α achieves the optimal ROC value according to:

$$\alpha^* = \operatorname{argmax}_{\alpha} \{ \text{TPR}_{\alpha} - \text{FPR}_{\alpha} \}$$

You can use `sklearn.metrics.roc_curve` function and the code provided in Lab 04. Using this value of α^* what is the model's test error?

10. Fit an ℓ_1 -regularized logistic regression by passing `penalty="l1"` when instantiating a logistic regression estimator
 - Set $\alpha = 0.5$
 - Use your previously implemented cross-validation procedure to choose λ
 - After selecting λ repeat fitting with the chosen λ and $\alpha = 0.5$ over the entire train portion.

What value of λ was selected and what is the model's test error?

11. Repeat question 10 for ℓ_2 regularized logistic regression. What value of λ was selected and what is the model's test error?